

Housing Prices

Ridge and Lasso regressions for the prediction of the median house price:
a machine learning implementation

Anna Olena Zhab'yak¹

University of Milan, Data science and Economics

and

Michele Maione²

University of Milan, Computer Science

(October 2020)

Key words: *machine learning, ridge regression, lasso regression, hedonic model*

Summary – This paper is trying to answer to the question what the best regression model is to predict the median house price. The hedonic theory is exploited and models as the Ridge and the Lasso regression are used on a cross sectional dataset of housing prices. Applying the cross-validation we estimate the parameters and then evaluate the model. Then PCA is implemented to improve the risk estimator. In section 1 we introduce the problem of housing prices in the U.S. and the approach used in this work. Section 2 is dedicated to the literature about regressions in predicting the price of real estates and 3rd section the theoretical notations are clarified to simplify the understanding of *the notions*. Our experiment is described in section 4 and in section 5 the consequential critical comments and evaluations.

market prices, so extending this concept to the housing prices we can see the attributes as the house's characteristics that are determinant for the final value. Hedonic model exploits the consumer theory and her willingness to pay depending on the utility gained from the bundle of aggregated attributes. Each attribute differently influences the price and its strength is given by the estimated coefficient.

Our work starts from a real problem of housing prices in the United States, where the economical purposes and the low mortgage rates incentive a solid and hot real estate market.³ Indeed, the U.S. is one of the most stable and secure countries for real estate investment in the recent years⁴. It is estimated that household wealth is nearly 50% invested in real estate and the owner-occupied

1. Introduction and description of the problem

The hedonic theory identifies the attributes as implicitly embodied in goods and their observable

¹ annaolena.zhab@studenti.unimi.it – 960298.

² michele.maione@studenti.unimi.it – 931468.

³ Santarelli, 2020.

⁴ Source: International Investor Survey.

housing rate in July 2019 was about 63.5%⁵. However, the U.S. real estate market was not always as reliable as today, indeed the sudden bubble of the housing market of 2006-2007 preceding the Great Recession and its subsequent burst is clear evidence of the system weaknesses. The speculation on the housing prices and their extremely high values is due to the lack of information caused by the manipulations of major players in the real estate sector⁶. For these reasons, the task of predicting the value of a house becomes crucial, as the constructed house price model can influence the economic growth and improve the efficiency of the real estate market. An accurate prediction model is significant and helps to fill up an information gap for the prospective homeowners, policy makers and other real estate market participants, such as, mortgage lenders and insurers⁷. Modelling house prices presents some issues, for example the median value might be extremely influenced by the value of the sold properties in the area with similar characteristics⁸ or the prediction could become wrong due to exogenous factors influencing the prices. Indeed, the economic health reflects in the market according to the supply and demand law so any shock will affect the current prices. Moreover, working on a large dataset, like the one used in this work, can lead to the so-called multicollinearity of the features which tend to overfit when it comes to implement the algorithm predicting the value. The classic OLS regression has the desired property of being unbiased, but it can suffer of overfitting and have a huge variance in those cases where features are highly correlated. To pull down the variance and obtain more biased estimator a regularization technique is necessary. The focus of this paper is therefore on two regularization techniques, the Ridge and Lasso regression. The Ridge regression⁹ is a useful tool for improving prediction in regression tasks with highly correlated predictors. Lasso regression is also used to handle high dimensional databases where the features are correlated, and this technique shrinks some of them to zero, performing a feature selection with a consequent dimension reduction. Both methods

act on the coefficients by introducing a penalty on them to make more effort to the most informative ones, this way minimizing overfitting of the data and solving the multicollinearity problem. The impact of each attribute on the predicted price is given by the value of the coefficient, higher coefficients mean higher influence. The penalty is the tool through which we perform the regularization, also called tuning parameter, it controls the bias-variance trade-off and the selection of it is crucial. For choosing the regularization parameter in practice, cross-validation (CV) is widely used.

2. Most important related works

Many works have been developed to predict the median house value with models of different complexity [see Manjula et al., 2017]. The concept of hedonic prices was developed by Rosen (1974), however the first the first implementing the hedonic model to the house sector was Lancaster (1966). Griliches (1971) provided the reading of a commodity, such as a house, as an aggregation of individual components or attributes. Timothy Oladunni & Sharad Sharma (2016) and Limsombunchai et al. (2004) have showed that the price of a property is predictable exploiting the hedonic theory, comparing the hedonic regression in comparison with other algorithms. Dubin (1998) has developed a work to predicted house prices using MLS data, even though exploiting different algorithms for the prediction, such as kriging algorithm to create an accurate spatial interpolation of house prices. Others as Xin and Khalid (2018) have used ridge and lasso regression to deal with multicollinearity of features on a time series database for predicting the housing price. Hoerl and Kennard (1970) firstly introduced the Ridge regression as biased estimator for non-orthogonal problems. The asymptotic properties of ridge have been widely studied, [see for e.g. Dobriban and Wager (2018), Dicker (2016)]. For the validation approach we refer to the cross-validation which biased estimation of the error is known (Hastie et al., 2009, p. 243), since it uses a smaller amount of data than

⁵ Source: United States Census Bureau.

⁶ Oladunni, Timothy & Sharma, Sharad, 2016.

⁷ Limsombunchai et al. (2004).

⁸ The so-called Sales comparison approach.

⁹ Introduced by Hoerl and Kennard (1970).

the entire dataset.¹⁰ However, we can apply a bias-control, see Liu and Dobriban (2020), for example via k-fold cross validation, see Ray (2018), since there is an inverse relation between the k size and bias, if the first grows the latter goes down.

Mishra et al. (2017) have clearly explained the intuition behind the PCA and the underlying algebra to rich these results. PCA was introduced by Pearson (1901) and Hotelling (1933) and it is largely used in a lot of fields. Gupta and Kabundi (2010) have implemented lasso, ridge and PCA to predict housing prices on a time series dataset.

3. Notation and relevant definitions - Regression

The goal of the regression is to generate a prediction $\hat{y} = f(w, x)$ such that the loss function $\ell(y, \hat{y})$ is small for most data points $x \in \mathcal{X}$, where $\hat{y} \in \mathcal{Y}$ is the prediction from the labels set $\mathcal{Y} \subseteq \mathbb{R}$, $w \in \mathbb{R}^d$ is the coefficient vector and $\mathcal{X} = \mathbb{R}^d$ the data domain; the prediction mistakes are a function of the difference $|y - \hat{y}|$.

3.1. Hedonic model

Following the hedonic theory, the housing price can be written as a function $f(\cdot)$ in the following way:

$$P_i = f(s) \quad (1)$$

where s is the vector of all the objective attributes and P_i is the price of the i^{th} element of the data matrix \mathcal{X} . In this case, the price (our target variable) is a function of:

- longitude,
- latitude,
- housingMedianAge,
- totalRooms,
- totalBedrooms,
- population,
- households,
- medianIncome,
- medianHouseValue,

- oceanProximity.

3.2. Loss function

With loss function we denote the measure of how different the prediction of a hypothesis is from the true outcome. We use a nonnegative loss function to measure the discrepancy $\ell(y, \hat{y})$ between the predicted label \hat{y} and the true label y . In the regression task we define the quadratic loss that is the squared distance between y and \hat{y}

$$\ell(y, \hat{y}) = (y - \hat{y})^2 \quad (2)$$

when $\hat{y} = y$ then $\ell(y, \hat{y}) = 0$ otherwise If $\hat{y} - y = c, \forall c \in \mathbb{R}^+$ and c is large then also $\ell(y, \hat{y})$ tend to be large. The mean of the squared error (MSE) will be used in the experiment.

3.3. Test error and training error

The split of dataset into two separate subsets is necessary in order to have some fresh data to estimate the predictive power of the algorithm. The dataset is divided in n elements for the test, and m elements for the training. Indeed, the validation is given by the test error which is:

$$\frac{1}{n} \sum_{t=1}^n \ell(y'_t, f(x'_t)) \quad (3)$$

The validation is done over a fitted predictor in the training set, and its power is given by the training error:

$$\hat{\ell}(f) = \frac{1}{m} \sum_{t=1}^m \ell(y_t, f(x_t)) \quad (4)$$

Total error is given by three elements:

- variance,
- bias,
- irreducible error.

The main idea is to derive a trade-off between the bias and variance, on order to optimize them both. More complex models present high variance and low bias since they fit good the true data but generalize worst.

¹⁰ In other words, the algorithm has not enough data to train and can be approximated.




Figure 1 - the optimal choose for trade-off. Source: re-searchgate.net

3.4. Empirical Risk Minimization (ERM)

The Empirical Risk Minimization is a learning algorithm which returns some predictors $f \in \mathcal{F}$ given a set of predictors, that minimize the training error, given a non-negative real-valued loss function $\hat{\ell}$:

$$\hat{f} \in \operatorname{argmin}_{f \in \mathcal{F}} \hat{\ell}(f) \quad (5)$$

3.5. Statistical risk, Bayes optimal predictor and Bayes optimal risk

We use statistical learning to introduce the notion of expectations in estimating the loss since we need to assume the independence between the variables and the predictor we generate is based on this assumption. Let $h : \mathcal{X} \rightarrow \mathcal{Y}$ be the predictor that maps data points to labels. The **statistical risk** is then defined as the expectation of the loss function among D , the distribution from where the random sample of data points and labels were drawn:

$$\ell_D(h) = \mathbb{E}[\ell(Y, h(x))] \quad (6)$$

where $h(x)$ is the predicted \hat{y} . We then define as **Bayes optimal predictor** as the function f^* which minimize the overall training error $\ell_D(h)$, given the conditional probability among all predictors given that our data point is x :

$$f^*(x) = \operatorname{argmin}_{\hat{y} \in \mathcal{Y}} \mathbb{E}[\ell(Y, \hat{y}) | X = x] \quad (7)$$

The **Bayes optimal risk** is the expectation over the loss function of the Bayes optimal predictor and following the same logic as before we have that the Bayes risk is smaller than the other risks:

$$\mathbb{E}[\ell(Y, f^*(x))] \leq \mathbb{E}[\ell(Y, h(x))] \quad (8)$$

Coming to our regression problem with the squared loss, the Bayes optimal predictor is:

$$f^*(x) = \operatorname{argmin}_{\hat{y} \in \mathcal{Y}} \mathbb{E}[(Y - \hat{y})^2 | X = x] \quad (9)$$

minimizing this quantity¹¹, we have:

$$f^*(x) = \mathbb{E}[Y | X = x] \quad (10)$$

and the Bayes risk becomes the expectation of (6):

$$\mathbb{E}[(Y - f^*(X))^2 | X = x] = \operatorname{Var}[Y | X = x] \quad (11)$$

3.6. Regressions – Linear, Ridge, Lasso

Ridge and Lasso regression modify the standard linear regression by introducing a positive constant as regularization parameter. Indeed, the objective function to minimize under these solutions is **RSS¹² + penalty**, and the penalty differs for the two methods. Starting from the classical linear model we have:

$$y_i = \bar{x}_i^T w \quad (12)$$

let be the data domain $\mathcal{X} = \mathbb{R}^d$ and $x = (1, x_1, \dots, x_d)^{13}$ a row vector of \mathcal{X} . The linear predictor is a linear function $h : \mathbb{R}^d \rightarrow \mathbb{R}$, and for an activation function $f : \mathbb{R} \rightarrow \mathbb{R}$ we can write as follows:

$$h(x) = f(w^T x) \quad (13)$$

where $w \in \mathbb{R}^d$ and $w^T x = \sum_{i=1}^d w_i x_i$.

The Bayes optimal risk is given by

$$f^*(x) = \mathbb{E}[y | X = x] \quad (14)$$

and it is also an empirical risk minimization to $(x_1, y_1) \dots (x_m, y_m)$ is

¹¹ By taking the derivative w.r.t \hat{y} , since $f^*(x)$ is differentiable.

¹² Sum of the squared residuals used for the classical OLS.

¹³ Add one extra feature to stabilize the prediction.

$$\hat{w} = \operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{m} \sum_{t=1}^m (w^\top x_t - y_t)^2 \quad (15)$$

Since we can rewrite these terms in vector notation, we have

$$\hat{w} = \operatorname{argmin}_{w \in \mathbb{R}^d} \|v - y\|^2 \quad (16)$$

for $v = (w^\top x_1, \dots, w^\top x_m)$ the vector of predictions and $y = (y_1, \dots, y_m)$ the vector of real labels and for $v, y \in \mathbb{R}^m$.

In matrix notation we have S the design matrix $S \in \mathbb{R}^{m \times d}$ with d features and m observations x_i that are rows of S^\top , and therefore the vector becomes $v = Sw$. Applying the ERM we derive

$$\hat{w} = \operatorname{argmin}_{w \in \mathbb{R}^d} \|Sw - y\|^2 \quad (17)$$

The solution to the ERM is the minimization of this convex function $F(w) = \|Sw - y\|^2$ using the Euclidean norm. To solve the problem in linear regression we can use the closed form solution, or the gradient descend.

If $S^\top S$ is a non-singular matrix¹⁴, and the conditions of the general position holds, the solution of the ERM is the closed form:

$$\nabla F(w) = 2S^\top(Sw - y) = 0 \quad (18)$$

$$\hat{w} = (S^\top S)^{-1} S^\top y \quad (19)$$

In some cases, the linear regression performs well on the training data, having a low bias, but it gives a non-accurate estimate on different data. The reason why it occurs it is because of multicollinearity of the prediction vectors (as known as non-orthogonality)¹⁵. More in general with d large or n small, the risk that the model can overfit¹⁶ the data is high. The OLS estimator \hat{w} therefore is unbiased but have a huge variance and it is not stable. To overcome this problem, Ridge and Lasso regression help to prevent over-fitting which results from simple linear regression. We introduce a regularized parameter α which adds some bias¹⁷

whereas pushing the variance down. This also controls the model complexity, indeed the value of α has a direct relation with the complexity. This occurs to find the best trade-off between bias and variance to get to that sweet spot for having good predictive performance¹⁸.

The two methods work similarly but lead to different results, this happens because of the divergent formulas.

3.6.1. Ridge solution

Ridge regression uses the penalty multiplied by the square of the magnitude of the coefficients, also known as L2 regularization.

The ERM functional of Ridge regression is

$$\hat{w}_\alpha = \operatorname{argmin}_{w \in \mathbb{R}^d} \|Sw - y\|^2 + \alpha \|w\|^2 : \forall \alpha > 0 \quad (20)$$

for $\alpha \rightarrow 0$, $\hat{w}_\alpha \rightarrow \hat{w}$ so the solution leads the linear regression, for $\alpha \rightarrow \infty$ the coefficient tend to a zero vector and the line becomes flatter, shrinking the linear regression solution towards to zero.

To optimize the objective function, we take the gradient as before and solve for w to find a suitable value

$$\nabla F(w) = \|Sw - y\|^2 + \alpha \|w\|^2 \quad (21)$$

$$2S^\top Sw - S^\top y + 2\alpha w = 0 \quad (22)$$

$$(S^\top S + \alpha I)w = S^\top y \quad (23)$$

The new estimated parameter becomes:

$$\hat{w}_\alpha = (S^\top S + \alpha I)^{-1} S^\top y \quad (24)$$

This is the so called closed-form solution and α is the one measuring the stability of the procedure.

3.6.2. Lasso solution

Least Absolute Shrinkage and Selection Operator, or simply Lasso, is slightly different from the pre-

¹⁴ This happens if the data points span $m \geq d$.

¹⁵ Hoerl and Kennard, 2010.

¹⁶ Overfitting: the algorithm performs very good on training data but cannot be generalized to a new bunch of data.

¹⁷ Bias is how well the fit correspond to the true value.

¹⁸ See graphic 1 in this paper.

¹⁹ Adding the identity matrix fixes the invertibility problem, always compute inverse, and this is more stable solution.

vious because the penalty is multiplied by the absolute value of the magnitude of coefficients, also known as L1 regularization

$$\hat{w}_{Lasso} = \operatorname{argmin}_{w \in \mathbb{R}^d} \|Sw - y\|^2 + \alpha|w| : \forall \alpha > 0 \quad (25)$$

For $\alpha \rightarrow \infty$, $\hat{w}_{Lasso} = 0$. The Lasso procedure encourages simple, sparse models²⁰, indeed some coefficients can become zero and be eliminated from the model, this way performing a feature selection.

The shrinkage amount is given by the value of tuning parameter α . If α increase, we have some parameters go straightway to zero.

The optimization of a non-differentiable function as Lasso solution is done by a proximal gradient descend approach²¹.

The first step is to take the gradient descend for current $w^{(k)}$ vector and form a new vector $z^{(k)}$:

$$z^{(k)} = w^{(k)} - \eta X^T(Xw^{(k)} - y) \quad (26)$$

Where η is the step size and k is the moment we are considering. Then solve the proximal regularize problem for $w^{(k+1)}$ as follows:

$$w^{(k+1)} = \operatorname{argmin}_{w \in \mathbb{R}^d} \|z^{(k)} - w\|_2^2 + \eta\alpha|w| : \alpha, \eta > 0 \quad (27)$$

This is a scalar minimization problem indeed we can rewrite it as:

$$\min_{w_i} \sum_{i=1}^m \left(z_i^{(k)} - w_i \right)^2 + \alpha\eta|w_i| \quad (28)$$

Since we have an absolute value for $|w_i|$ we consider two cases:

Case 1 $w_i \geq 0$

$$\min_{w_i} \left(z_i^{(k)} - w_i \right)^2 + \alpha\eta w_i \quad (29)$$

differentiate with respect to w and solve:

$$-2(z_i - w_i)^2 + \alpha\eta = 0 \quad (30)$$

since we have the non-negativity constraint over w_i :

$$w_i = \begin{cases} z_i - \frac{\alpha\eta}{2}, & z_i > \frac{\alpha\eta}{2} \\ 0, & \text{else} \end{cases} \quad (31)$$

Case 2 $w_i < 0$

$$\min_{w_i} \left(z_i^{(k)} - w_i \right)^2 - \alpha\eta w_i \quad (32)$$

differentiate with respect to w and solve for it:

$$-2(z_i - w_i)^2 - \alpha\eta = 0 \quad (33)$$

$$w_i = \begin{cases} 0, & z_i + \frac{\alpha\eta}{2} > 0 \\ z_i + \frac{\alpha\eta}{2}, & \text{else} \end{cases} \quad (34)$$

This solution is also known as the “soft threshold” operation. More compactly we can rewrite the values for w_i as follows:

$$w_i = \begin{cases} 0, & \frac{-\alpha\eta}{2} < z_i < \frac{\alpha\eta}{2} \\ z_i - \frac{\alpha\eta}{2}, & z_i > \frac{\alpha\eta}{2} \\ z_i + \frac{\alpha\eta}{2}, & z_i < \frac{-\alpha\eta}{2} \end{cases} \quad (35)$$

The common point of these two methods is that adding the regularization parameter to the cost function the algorithm is forced to pick the lowest weights, indeed the goal is to ensure a small coefficient through this regularization parameter. The main difference is that many coefficients are exactly zeroed under lasso, which is never the case in ridge regression where there is not any elimination of coefficients. Moreover, Lasso arbitrarily selects any one feature among the highly correlated ones, leading to a higher variance than Ridge regression.

3.7. Cross-validation

The performance of an algorithm has to be evaluated on a new data order to see if it works even on a not trained data. The most common approach used is to split the dataset into three subsets,

²⁰ Stephanie Glen. “Lasso Regression: Simple Definition” from “StatisticsHowTo.com: Elementary Statistics for the rest of us!”.

²¹ This method requires a convex function. Lasso lies in this category.

training, validation set and test set. The first set will be used to train the algorithm and the second to evaluate its performance, finally the best output according to the validation set will be evaluated on the final test set to estimate the risk.

Cross-validation (CV) is one of the techniques used to test the effectiveness of a machine learning model and it is also a re-sampling procedure used to evaluate a model if we have a limited data. The CV risk estimate therefore becomes $\mathbb{E}[\ell_D(A(s))]$, where $A(s) = \hat{h}$ is our estimated predictor on the training set. This approach however can be biased²², therefore a K-Fold cross-validation is largely used for evaluating the accuracy of model.

This approach splits the training set into K-subsets and each fold at each interaction is used for testing while the remaining are used to training. This ensures that every observation from the original dataset has the chance of appearing in training and test set, this way decreasing the bias of the CV²³.

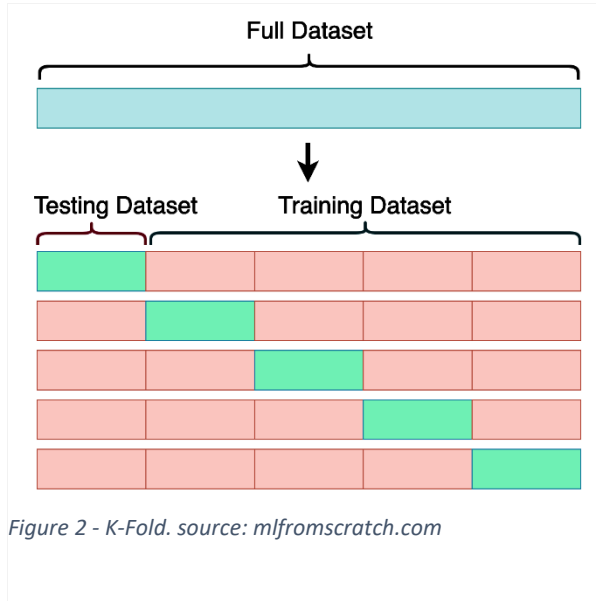


Figure 2 - K-Fold. source: mlfromscratch.com

We have D_k the testing part while $S^{(k)}$ the train part. The number of subsets obtained is k and we get a value for $h_k = A(S^{(k)})$ computed on the training part. The test loss according to k-fold is

$$\hat{\ell}_{D_k}(h_k) = \frac{K}{m} \sum_{(x,y) \in D_k} \ell(y, h_k(x)) \quad (36)$$

After all interactions (from 1 until k^{th} interaction) we collect the obtained values, and the CV risk estimate is the average of all errors:

$$\frac{1}{K} \sum_{k=1}^K \hat{\ell}_{D_k}(h_k) \quad (37)$$

The complete algorithm works like this:

1. Shuffle the dataset randomly,
2. Split the dataset into k groups,
3. For each unique group:
 1. Take the group as a hold out or test data set,
 2. Take the remaining groups as a training data set,
 3. Fit a model on the training set and evaluate it on the test set,
 4. Retain the evaluation score and discard the model.
4. Summarize the skill of the model using the sample of model evaluation scores.

However, this approach cannot be used to tune the hyperparameter that is involved into the estimation of the parameter such that to obtain $\min_{\theta \in \Theta} \ell_D(h_s^{(\theta)})$ where $h_s^{(\theta)} = A_\theta(S)$. The choice of it must be done before the choice of the training set on which the algorithm will be learned. To do this we use another variation of CV that is the Nested Cross-validation, which puts together the previous two techniques in such a way to solve the hyperparameter choice problem. The set is split into k-fold, one for testing and the rest for training. After two loops are made:

²² Because we could be just lucky with the train set we had and the risk estimated on the test set suffer of this distortion.

²³ As $k \rightarrow N$, leads to "Leave-one-out cross-validation", the leave one out approach, where the number of sets equal the number of observations. On the contrary for smaller values of k we have the CV approach.

- An internal cross validation on the training part of the initial split is also split in k-folds and run the validation for a grid of all $\theta \in \Theta$ choosing the best one.
- External cross validation is done in order to avoid the dependency over the training set, and it takes the best output from the inner loop and run it on the entire training set. The risk is therefore evaluated on the remaining test set of the current external fold.

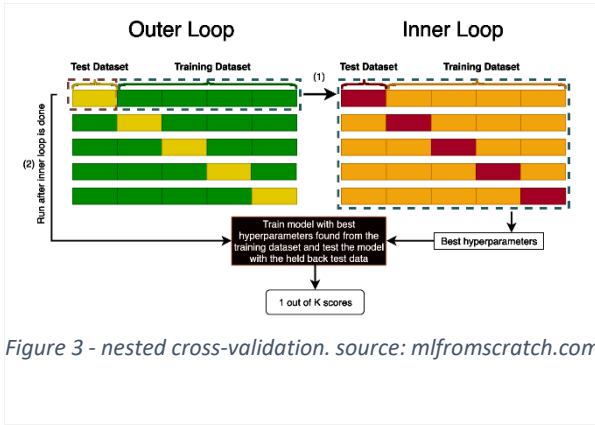


Figure 3 - nested cross-validation. source: mlfromscratch.com

This procedure is run on each fold and at the end we pick up the expected value of the average risk as follows

$$\mathbb{E} \left[\min_{\theta \in \Theta} \ell_D(A_\theta) \right] \quad (38)$$

3.8. Principal component analysis

Principal Component Analysis is a technique used for dimensionality reduction. Its goal is to reduce the number of features through a combination of the original data variables, in this way keeping most of the original information. Standardization is needed. This technique finds the eigenvalues and eigenvectors of the correlation matrix²⁴.

The selection of the principal components is based on the variance caused to the target variable. The principal components will be then independent one from the other. The feature that cause more variance is the first principal component, and so

on until we reach a suitable number of explained variance by the principal components.

A good way is to plot the variance against principal components and ignore the principal components with diminishing values as shown in the following graph:²⁵ This way we reduce the variance and we can improve the stability of the regression by solving the multicollinearity. It is a way of identifying patterns in data and expressing the data in such a way as to highlight their similarities and differences.

4. Proof of a technical result

The demonstration of our experiment and all the critical considerations will be described here. This material is also available on GitHub at this url: <https://github.com/mikymaione/HousingPrices> we have created a Jupiter Notebook to illustrate the procedure followed step by step at this url: <https://github.com/mikymaione/HousingPrices/blob/master/SourceCode/HousingPrices/main.ipynb>

4.1. Data pre-processing

Before performing the analysis and regression, pre-processing of data is necessary. The dataset presents features that cannot be compared in a linear Euclidian space; therefore, geometry is not working properly on this row data. Indeed, to learn the algorithm, we need to encode the features and raise them to a homogeneous level, so we can compare them²⁶. The dataset contains 20,640 observations and 10 features for each house including the median house value which is the target value that we are trying to predict. Firstly, we create the two constants of the target variable and designed matrix: X , y .

4.1.1. Missing values

The missing values must be handled to avoid errors in the execution of the code, so they are filled with the mean value of the corresponding column.

²⁴ The eigenvectors of the Covariance matrix are actually the directions of the axes where there is the most variance (most information).

²⁵ Usman Malik.

²⁶ For example, longitude and households are features with real numbers however they have different interpretations.

4.1.2. Categorical features

There is a categorical feature which represents the distance from the ocean, we transform the elements of the column into columns dummies and assign it to the data set²⁷. Even though among hedonic model literature²⁸ there is a concern about the statistical insignificance of some features such as household size, we use two approaches:

- keep all the features in regressions to have as much information as possible,
- use some unsupervised techniques to decide which feature to drop such as PCA.

4.1.3. Standardization

The standardization is done by subtracting the mean and divided by the variance. In this way we have $\mu = 0$ and $\sigma = 1$. This procedure is needed as both L1 and L2 assume that all features are centred around 0 and have unit variance. More in general this is important to those algorithms that use Euclidian distance and for PCA implementation.

4.1.4. Correlation matrix

We explore the correlation matrix between the features to see graphically how they move together. The used scoring is Pearson's coefficient r and $r \in [-1; 1]$, if the value is closer to 1 there is more correlation and the sign gives the direction of this correlation. Darker and lighter colours on this map are the two extremes.

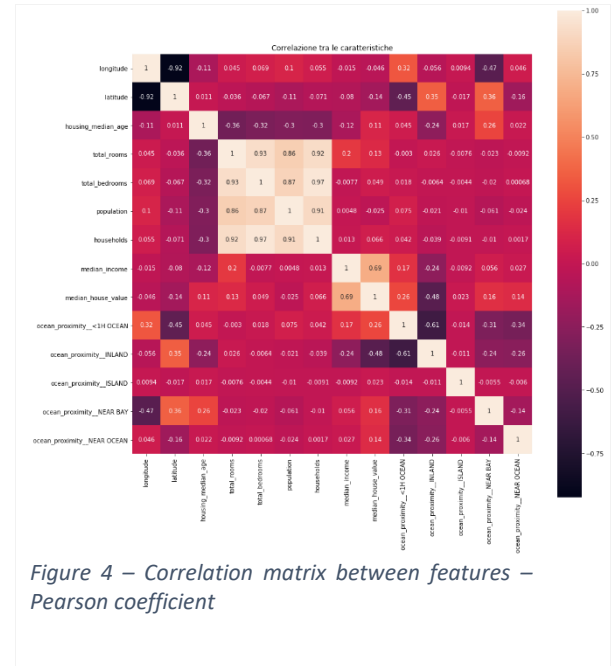


Figure 4 – Correlation matrix between features – Pearson coefficient

We have enough evidence that there exists statistical relationship between the variables²⁹. Therefore, our dataset is suitable for decomposition into its principal components to increase convergence speed and eliminate collinearity by finding the core components of the datasets.

4.2. Model tuning

Model tuning consists in the choice of the parameters to use into regression. In our case is the α hyperparameter for both Ridge and Lasso regression. The train-test split is done 80% - 20%³⁰

4.2.1. Scoring

Mean squared error is implemented as scoring, and it takes bigger values more than proportionally if the error in prediction increases.

4.2.2. Choosing the set of parameters alpha

In order to obtain a reasonable amount of information to determine a certain $f^* \in \mathcal{F}$ where f^* is the function that minimize the training error, we need to set different values of the tuning parameter to find out the best one. A larger value of α leads to a high bias but a low variance. On the

²⁷ This is called hot encoder technique; we do not worry about adding extra dimensions as the dummy variable sets to zero the features that do not belong to the given observation.

²⁸ Features such as garden, household size, neighborhood satisfaction and schools, as demonstrated by Berna and Craig.

²⁹ See the lighter square 4x4 in the middle.

³⁰ 80/20 rule: following the Pareto principle.

other hand, for small values of α the variance increases, and bias go down. We perform the analysis on α with the relative mean squared error on the training data, comparing the Lasso and Ridge solutions. Nested cross validation gives us the same best value as the non-nested cross validation. We are plotting the parameter using the non-nested CV and then comparing the two methods.

Ridge regression

We use a logarithmic range $\ln|F|$. Training size m is bigger than the $\ln|F|$ in order to avoid underfitting.




Figure 5 - Ridge regression: validation curve with $\alpha \in (0; 0.8)$

As we can see the optimal value for the hyper-parameter that optimize the squared loss is in between the range $(0; 0.1]$, after that the squared loss increases. The best value of the penalized term is **0.000059**.




Figure 6 - Ridge: nested cross-validation vs non-nested cross-validation of the best α

Lasso regression

For Lasso implementation we have these values:

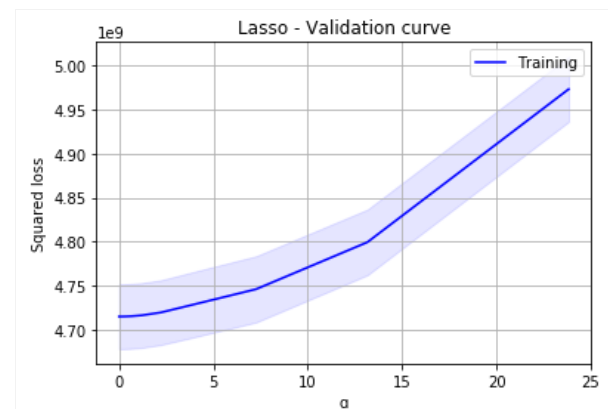


Figure 7 – Lasso regression: validation curve with $\alpha \in (0; 24)$

The values of alpha used are linear, and the best alpha according to the validation curve is **0.0003**.




Figure 8 – Lasso: nested cross-validation vs non-nested cross-validation of the best α

4.3. Learning algorithm

Once we have determined the hyperparameter, the optimization of the learning algorithm is done. For ridge regression we use the Cholesky method, that is the closed form. For lasso regression we apply the proximate gradient descend³¹ This procedure is internal to the ridge and lasso function, and the learning algorithm is the output.

4.3.1. Ridge learning algorithm

We fit the best α to plot the learning curve performance. Training error becomes larger when iterations are increased, and test error is higher as we could always expect a better performance on the training set. As we see the overfitting disappears as we increase the training size, around 5,000 training size, and it is improving with training size growth with a stable squared loss (4.7 in Figure 9).




Figure 9 - Ridge regression: learning curve with different training set sizes

In this plot we can visualize how the predicted value differs from the real values. This is done with fit and predict functions. The prediction is more consistent with lower prices and becomes sparser for higher values, this can be caused by the presence of the outliers. The R^2 is around 63%.




Figure 10 - Ridge regression: scatter plot prediction vs test

Here we can see the magnitude of each coefficient, and its prediction power on the target variable. The overall picture is very similar to the ridge coefficient, island location is the driven feature of the housing price.

³¹ This is possible because we take an approximation of the gradient since Lasso function is not derivable.




Figure 11 - Ridge: coefficients magnitude

Looking at the coefficients of ridge regression we can conclude that household size is statistically insignificant as the literature suggests. On the other hand, the driven force is the house location in island, which presents a direct and positive impact on the predicted prices. Also, the median income causes higher prices. On the other hand the Inland location have a negative impact on the prices.

4.3.2. Lasso learning algorithm

In lasso regression this increase is smoother, and MSE is more stable for α in between (0; 1]. This is what we could expect if there are feature highly correlated and that are not crucial for the regression. On the other hand, R^2 is specular to the MSE and tells u show well the model fits the data.

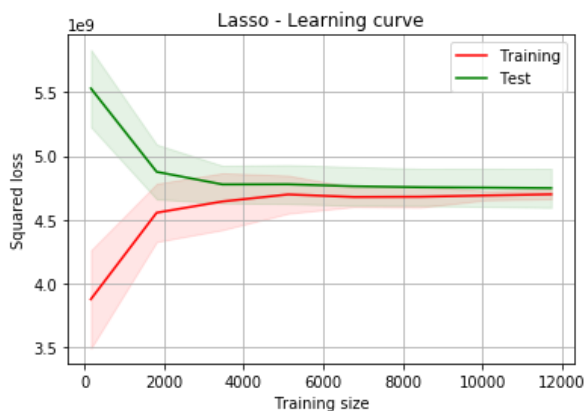


Figure 12 - Lasso regression: learning curve with different training set sizes

The coefficient magnitude is shown in this graph.




Figure 13 - Lasso: coefficients magnitude

Median income has a huge power in the prediction, it means that it drives the values. They are very similar to the coefficients of Ridge regression (**Error! Reference source not found.**) however some values differs by a very small amount.

4.4. Principal Component Analysis

Cells that are highly correlated cluster together. Differences between the 1st pc (plotted on x axes) are more important than the differences between the 2nd pc. Each point is a predictor that we have learned. Now we have a look on the Principal components of this data set and how is the variance distributed among two principal components.




Figure 14 - PCA for 2 principal components on the axes. In the graph we have the predictors and their contribution to the variance.

As we can see there are two main predictors that cause the first largest variance on x-axes and the second variation on the y-axes. The negative values mean the existence of an inverse correlation between the factor PCA and the variables.




Figure 15 - Plot of the learned model after PCA

The group of points that are close to each other are more similar and do not have that much impact on the two principal components. In the next plot we see how the cumulative variance is explained by each feature (in our case 13).

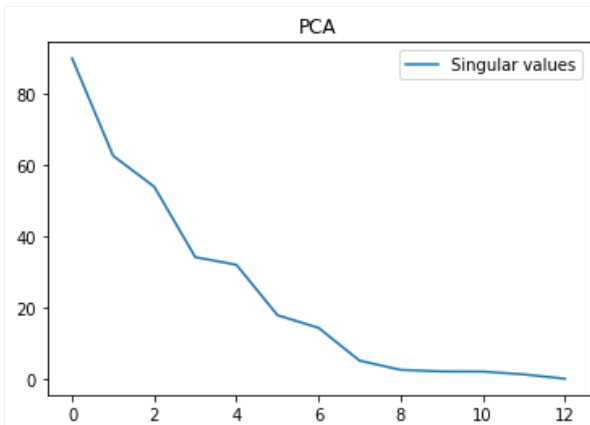


Figure 16 - PCA for singular values vs cumulative variance

After 4 features the variance significantly drops and above 7 features we do not gain more information, indeed the variance is almost 0. We will implement the decomposition on 7 features.

This is the result and the performance of 7 PCA decomposition using our learned variables and projecting the datapoints.




Figure 17 - PCA: test and train loss for the 7-PCA decomposition

The performance of the curve is similar to the one of the Ridge regression in Figure 12, however the variance is still high. In the following graph we plot again the principal components using the learned variables.

As we can see the spread is not improved that much but now we have just one variable that causes the first large variance. The second main variance is due to another variable meanwhile the remaining have less spread. This could be maybe the inland location of the house and the median income as it resulted from ridge regression coefficients in Figure 11.

5. Some critical considerations

We have performed two different regression on the dataset of houses, and it results in a very similar performance for both Ridge and Lasso regularizations, this can be because in both methods we are using the best nested cross-validated penalized parameter.




Figure 18 - Lasso: scatter plot prediction vs test

Moreover, we have tried to improve the risk using the PCA, but the evidence shows pretty similar results as well. This can be because both PCA and Ridge regression are methods are designed to have a stabilized regression.

In general, for this specific database we have noticed if we drop some feature or elements such as outliers the predictive power is poor. Indeed, the performance is better when all the features are considered in the prediction as we have done in this experiment.

6. Bibliographical references

- Berna and Creig, 2016. "Defining spatial housing submarkets: Exploring the case for expert delineated boundaries". SAGE journals.
- Calhoun C. A., 2003, "Property Valuation Models and House Price Indexes for The Provinces of Thailand: 1992 –2000", Housing Finance International, 17(3): 31–41
- Dicker, Lee H. Ridge regression and asymptotic minimax estimation over spheres of growing dimension. Bernoulli 22 (2016), no. 1, 1--37. <https://projecteuclid.org/euclid.bj/1443620842>
- Dobriban, Edgar & Wager, Stefan. (2015). High-Dimensional Asymptotics of Prediction: Ridge Regression and Classification. The Annals of Statistics. 46
- Dubin, Robin. (1998). Predicting House Prices Using Multiple Listings Data. The Journal of Real Estate Finance and Economics. 17, 35-59
- Frew J. and B. Wilson, 2000, "Estimation The Connection Between Location and Property Value", Essay in Honor of James A.Graaskamp, Boston, MA: Kluwer Academic Publishers
- Frew J. and B. Wilson, 2000, "Estimation The Connection Between Location and Property Value", Essay in Honor of James A.Graaskamp, Boston, MA: Kluwer Academic Publishers
- Gupta, Kabundi, 2010, "Forecasting Real U.S.House Prices: Principal Components Versus Bayesian Regressions". International Business & Economics Research Journal.
- Hotelling H. 1933. Analysis of a complex of statistical variables into principal components. J Educ Psychol. 25: 417-441.
- Limsombunchai, Visit & Gan, Christopher & Lee, Minsoo. (2004). House Price Prediction: Hedonic Price Model vs. Artificial Neural Network. American Journal of Applied Sciences
- Liu, Sifan & Dobriban, Edgar. (2020). Ridge Regression: Structure, Cross-Validation, and Sketching
- Manjula, R & Jain, Shubham & Srivastava, Sharad & Kher, Pranav. (2017). Real estate value prediction using multivariate regression models. IOP Conference Series: Materials Science and Engineering. 263. 042098
- Mishra, Sidharth & Sarkar, Uttam & Taraphder, Subhash & Datta, Sanjoy & Swain, Devi & Saikhom, Reshma & Panda, Sasmita & Laishram, Menalsh. (2017). Principal Component Analysis. International Journal of Livestock Research.
- Oladunni, Timothy & Sharma, Sharad. (2016). Hedonic Housing Theory – A Machine Learning Investigation
- Pearson K. (1901) On lines and planes of closest fit to systems of points in space, Philosophical Magazine 2(11):559-572.
- Ray, 2018 www.analyticsvidhya.com/blog/2018/05/improve-model-performance-cross-validation-in-python-r
- Ridge Regression: Biased Estimation for Nonorthogonal Problems Author(s): Arthur E. Hoerl and Robert W. Kennard Source: Technometrics, Vol. 12, No. 1 (Feb., 1970), pp. 55-67 Published by: American Statistical Association and American Society for Quality Stable URL:

- Rosen S., 1974, "Hedonic Prices and Implicit Markets: Product Differentiation in Pure Competition", Journal of Political Economics, 82: 34 – 55
- Santarelli, 2020. "US Housing Market Forecast 2020 & 2021: Crash or Boom?". Norada Real Estate.
- Xin, Seng & Khalid, Kamil. (2018). Modeling House Price Using Ridge Regression and Lasso Regression. International Journal of Engineering & Technology. 7. 498
- Zvi Griliches, 1991. "Hedonic Price Indexes and the Measurement of Capital and Productivity: Some Historical Reflections," NBER Chapters, in: Fifty Years of Economic Measurement: The Jubilee of the Conference on Research in Income and Wealth, pages 185-206, National Bureau of Economic Research, Inc .

7. Sitographical references

- afire.org
- builtin.com
- census.gov
- datacamp.com
- jstor.org
- noradarealestate.com
- psu.edu
- researchgate.net
- stackabuse.com
- statisticsshowto.com
- towardsdatascience.com
- mlfromscratch.com

8. Table of figure

Figure 1 - the optimal choose for trade-off. Source: researchgate.net 2

Figure 2 - K-Fold. source: mlfromscratch.com 2

Figure 3 - nested cross-validation. source: mlfromscratch.com 2

Figure 4 – Correlation matrix between features – Pearson coefficient 2

Figure 5 - Ridge regression: validation curve with $\alpha \in (0; 0.8)$ 2

Figure 6 - Ridge: nested cross-validation vs non-nested cross-validation of the best α 2

Figure 7 – Lasso regression: validation curve with $\alpha \in (0; 24)$ 2

Figure 8 – Lasso: nested cross-validation vs non-nested cross-validation of the best α 2

Figure 9 - Ridge regression: learning curve with different training set sizes 2

Figure 10 - Ridge regression: scatter plot prediction vs test 2

Figure 11 - Ridge: coefficients magnitude 2

Figure 12 - Lasso regression: learning curve with different training set sizes 2

Figure 13 - Lasso: coefficients magnitude 2

Figure 14 – PCA for 2 principal components on the axes. In the graph we have the predictors and their contribution to the variance..... 2

Figure 15 - Plot of the learned model after PCA.. 2

Figure 16 - PCA for singular values vs cumulative variance..... 2

Figure 17 - PCA: test and train loss for the 7-PCA decomposition 2

Figure 18 - Lasso: scatter plot prediction vs test . 2

9. Copyright

We declare that this material, which We now submit for assessment, is entirely our own work and has not been taken from the work of others, save and to the extent that such work has been cited and acknowledged within the text of our work. We understand that plagiarism, collusion, and copying are grave and serious offences in the university and accept the penalties that would be imposed should I engage in plagiarism, collusion or copying. This assignment, or any part of it, has not been previously submitted by us or any other person for assessment on this or any other course of study.

9.1. MIT License

Copyright (c) 2020 Anna Olena Zhab'yak, Michele Maione

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and

associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

The software is provided "as is", without warranty of any kind, express or implied, including but not limited to the warranties of merchantability, fitness for a particular purpose and noninfringement. in no event shall the authors or copyright holders be liable for any claim, damages or other liability, whether in an action of contract, tort or otherwise, arising from, out of or in connection with the software or the use or other dealings in the software.